

# Decoding the Musical Genome

Lakshya Dugar

6/18/2020

## Contents

<b>PART 1 - Predicting Genres Based on Musical Genome</b>	<b>1</b>
INTRODUCTION . . . . .	1
LOADING LIBRARIES . . . . .	2
DATA WRANGLING . . . . .	3
FUNCTIONS and METHODS . . . . .	5
PILOT RESULTS . . . . .	8
ANALYSIS . . . . .	10
TRYING TO IMPROVE ACCURACY . . . . .	14
CONCLUSION . . . . .	16

## PART 1 - Predicting Genres Based on Musical Genome

### INTRODUCTION

Google defines music as a collection of coordinated sounds. Making music is said to be the process of putting sounds and tones in an order.

In the first part of this project I will try to determine if music can be decoded and if every genre can be explained and appropriately predicted in its simplest form, using the following data-points:

**Danceability** - the ease with which a person could dance to a song over the course of the whole song,  
**Energy** - how fast paced vs slow paced the song is,  
**Key** - the major or minor scale around which a piece of music revolves,  
**Loudness** - attribute of auditory sensation in terms of which sounds can be ordered on a scale extending from quiet to loud,  
**Mode** - a type of musical scale coupled with a set of characteristic melodic behaviors,  
**Speechiness** - the presence of spoken words in a track,  
**Acousticness** - describes how acoustic a song is,  
**Instrumentalness** - the amount of vocals in the song,  
**Liveness** - probability that the song was recorded with a live audience,  
**Valence** - the musical positiveness conveyed by a track,  
**Tempo** - the pace or speed at which a section of music is played (BPM),  
**Duration\_ms** - duration of the song in minutes,  
**time\_signature** - Release date.

We will try to predict the genres just from these criterion and see if genres are essentially just their genome or is there something else. Additionally, since this data set of just 131,580 songs has 626 genres, we will see what is the state of the overlap and if just using the genome is viable to predict the genres.

Aim - To see analysing just the musical genome is a viable model for the predictive analysis of the genres.

## LOADING LIBRARIES

```
if(!require(dplyr)) install.packages("dplyr", repos = "http://cran.us.r-project.org")
if(!require(caTools)) install.packages("caTools", repos = "http://cran.us.r-project.org")
if(!require(caret)) install.packages("caret", repos = "http://cran.us.r-project.org")
if(!require(class)) install.packages("class", repos = "http://cran.us.r-project.org")
if(!require(mlbench)) install.packages("mlbench", repos = "http://cran.us.r-project.org")
if(!require(glmnet)) install.packages("glmnet", repos = "http://cran.us.r-project.org")
if(!require(nnet)) install.packages("nnet", repos = "http://cran.us.r-project.org")
if(!require(DescTools)) install.packages("DescTools", repos = "http://cran.us.r-project.org")
if(!require(rpart)) install.packages("rpart", repos = "http://cran.us.r-project.org")
```

## DATA WRANGLING

To complete this classification exercise, we're going to borrow **Adri Molina's dataset** from Kaggle which contains a TSV file full of songs and features that will help us categorize the songs into groups (like time signature , key, and tempo).

Let's be sure to eliminate any columns that aren't useful features, and change any factors (besides the predicted factor, "Genre") to numerics, to simplify things

```
#Read in the file, which is tab-delimited
data.full <-
  read.delim("songDb.tsv", header = TRUE, sep = "\t")

# Remove columns that don't serve as features (the Spotify URI,
# The track reference, the full URL, etc.)
dataax <-
  subset(data.full,
         select = -c(Uri, Ref_Track, URL_features, Type, ID, Name))

# Identify each song by its name (by changing the row names to song names)
rownames(dataax) <- make.names(data.full$name, unique = TRUE)
# Ensure the time signature is numeric, rather than a factor
dataax$time_signature <- as.numeric(dataax$time_signature)

# Tempo should also be numeric
dataax$Tempo <- as.numeric(dataax$Tempo)
```

Now the data is ready in the format that is suitable for our analysis and we can proceed with the project.

This is what our data set looks like:

```
as_tibble(dataax)

## # A tibble: 131,580 x 14
##   Danceability Energy Key Loudness Mode Speechiness Acousticness
##       <dbl>    <dbl> <dbl>    <dbl> <dbl>      <dbl>        <dbl>
## 1     0.624    0.857    10     -6.25     0     0.0542     0.0208
## 2     0.517    0.916     0     -4.93     1     0.0559   0.000182
## 3     0.251    0.894     8     -4.10     0     0.057    0.0144
## 4     0.469    0.743     1     -5.57     0     0.0272   0.00222
## 5     0.487    0.952     1     -4.43     0     0.0613   0.000228
## 6     0.43     0.797     2     -5.91     0     0.0303   0.000308
## 7     0.434    0.908     6     -4.72     1     0.0936   0.00791
## 8     0.308    0.965     8     -3.17     1     0.0591   0.0000228
## 9     0.5     0.925     4     -3.47     0     0.0378   0.00094
## 10    0.479    0.977     2     -4.51     1     0.086    0.0000166
## # ... with 131,570 more rows, and 7 more variables: Instrumentalness <dbl>,
## #   Liveness <dbl>, Valence <dbl>, Tempo <dbl>, Duration_ms <dbl>,
## #   time_signature <dbl>, Genre <fct>
```

These are the summary statistics:

```
summary(datax)
```

```
##   Danceability      Energy       Key      Loudness
## Min. :0.0000  Min. : 0.0000  Min. :-14.372  Min. :-60.000
## 1st Qu.:0.4320 1st Qu.: 0.4870 1st Qu.: 2.000  1st Qu.:-10.377
## Median :0.5660 Median : 0.6900 Median : 5.000  Median : -7.377
## Mean   :0.5538 Mean  : 0.6488 Mean  : 5.311  Mean  : -8.523
## 3rd Qu.:0.6920 3rd Qu.: 0.8530 3rd Qu.: 9.000  3rd Qu.: -5.344
## Max.  :0.9880  Max. :11.0000 Max. : 11.000  Max. : 5.056
##
##        Mode      Speechiness     Acousticness Instrumentalness
## Min. :0.000  Min. :0.000000  Min. :0.000000  Min. :0.00000000
## 1st Qu.:0.000 1st Qu.:0.03590 1st Qu.:0.00793 1st Qu.:0.0000019
## Median :1.000 Median : 0.04830 Median : 0.10500 Median :0.0014800
## Mean   :0.619 Mean  : 0.08374 Mean  : 0.27099 Mean  :0.2318704
## 3rd Qu.:1.000 3rd Qu.:0.08300 3rd Qu.:0.48700 3rd Qu.:0.5030000
## Max.  :1.000  Max. :0.96600  Max. :0.99600  Max. :0.9990000
##
##        Liveness      Valence      Tempo Duration_ms
## Min. :0.0000  Min. : 0.0000  Min. : 1  Min. :     3
## 1st Qu.:0.0951 1st Qu.: 0.2470 1st Qu.:11846 1st Qu.: 190933
## Median :0.1250 Median : 0.4590 Median :22406 Median : 229000
## Mean   :0.1933 Mean  : 0.4938 Mean  :25051 Mean  : 253666
## 3rd Qu.:0.2460 3rd Qu.: 0.6830 3rd Qu.:39009 3rd Qu.: 285479
## Max.  :1.0000  Max. :187.8270 Max. :54265 Max. :5949886
##
##        time_signature          Genre
## Min.   : 1.000 alternativeamericana: 1891
## 1st Qu.: 4.000 electrolatino      : 1009
## Median : 4.000 doo-wop           : 972
## Mean   : 3.914 reading          : 969
## 3rd Qu.: 4.000 nuelectro        : 909
## Max.  :14.000 groovemetal      : 903
##                      (Other)        :124927
```

## FUNCTIONS and METHODS

Let's load the functions that we'll need for the first part of our project

### Test and Train Sets

Let's also define a function that makes it easy to create our train and test sets. (Storing both, the full train and test sets as well as the separate X and Y data frames for each might seem redundant but it will make things just a bit easier down the road)

```
get_train_test <- function(split_ratio, data) {  
  results <- list()  
  
  split.index <- sample.split(seq_len(nrow(data)), split_ratio)  
  
  results$data.train <- data[split.index, ]  
  results$data.test <- data[!split.index, ]  
  
  results$X.train <-  
    results$data.train %>% select(-Genre) %>% as.matrix()  
  results$Y.train <- results$data.train$Genre  
  
  results$X.test <-  
    results$data.test %>% select(-Genre) %>% as.matrix()  
  results$Y.test <- results$data.test$Genre  
  return(results)  
}
```

## kNN Model

Let's make a function which allows us to subset the entire dataset to contain songs only of a particular genre and will give us the accuracy of our algorithm using the k-Nearest-Neighbours model. Why we subset our data to include only specific genres will become clear later.

```
knn_function <- function(data, genres) {  
  data.sub <- data[data$Genre %in% genres, ]  
  data.sub$Genre <- droplevels(data.sub$Genre)  
  
  set.seed(101)  
  # Create an empty data frame to store the predictions and the actual labels  
  classifications <- data.frame(pred = factor(), actual = factor())  
  # Use K-fold cross validation  
  K = 5  
  for (k in 1:K) {  
    # shuffle the data  
    res <- get_train_test(0.8, data.sub)  
    fit.knn <-  
      knn(  
        train = res$X.train,  
        test = res$X.test,  
        cl = res$Y.train  
      )  
    classifications <-  
      rbind(classifications,  
            data.frame(pred = fit.knn, actual = res$Y.test))  
  }  
  confusionMatrix(classifications$pred, classifications$actual)  
}
```

## Decision Trees Model

Similar to the kNN, let's make a function which allows us the subset the entire dataset to contain songs only of a particular genre which will give us the accuracy of our algorithm using the decision trees model. Why we subset our data to include only specific genres will become clear later.

```
dtree_function <- function(data, genres) {  
  data.sub <- data[data$Genre %in% genres,]  
  data.sub$Genre <- droplevels(data.sub$Genre)  
  res <- get_train_test(0.8, data.sub)  
  
  # Decision Tree  
  set.seed(103)  
  fit.dtree <-  
    train(  
      Genre ~ .,  
      data = res$data.train,  
      method = "rpart",  
      parms = list(split = "information")  
    )  
  
  Y.pred.dtree <-  
    predict(fit.dtree, newdata = data.frame(res$X.test), type = "raw")  
  confusionMatrix(Y.pred.dtree, res$Y.test)  
}
```

## PILOT RESULTS

Let's sample 10 genres at random from the dataset and see what our accuracy is

```
set.seed(3)
genres <- sample(levels(datax$Genre), 10)
knn_function(datax, genres)$overall["Accuracy"]

## Accuracy
## 0.2753036

dtree_function(datax, genres)$overall["Accuracy"]

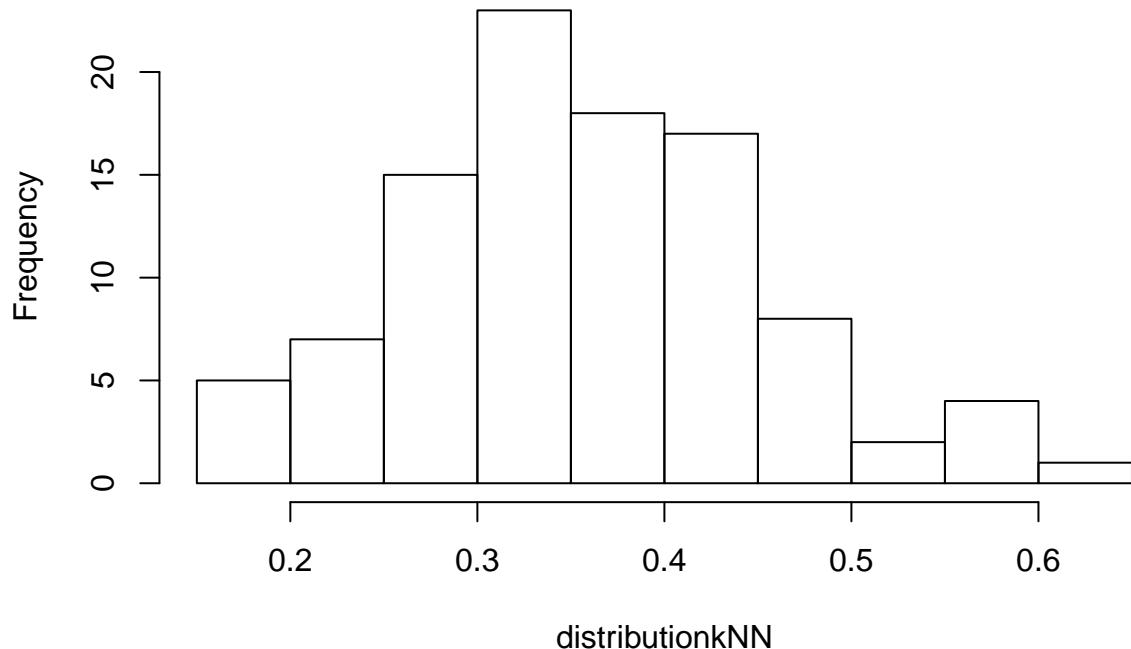
## Accuracy
## 0.3076923
```

Here's the histogram of the frequency distribution of the accuracy

```
set.seed(1)
distributionkNN <- replicate(100,{
  genres <- sample(levels(datax$Genre), 10)
  knn_function(datax, genres)$overall["Accuracy"]
})

hist(distributionkNN)
```

**Histogram of distributionkNN**



As we can see, a vast majority of the accuracy distribution is less than 0.5. Let's select 10 genres that we know to be sufficiently different from each other in terms of their genome composition and see if that increases the accuracy much.

```
genres <-  
  list(  
    "canadianpop",  
    "electronica",  
    "rock",  
    "modernblues",  
    "r&b",  
    "polishblackmetal",  
    "videogamemusic",  
    "irishfolk",  
    "koreanpop",  
    "hiphop"  
)  
knn_function(datax, genres)$overall["Accuracy"]  
  
## Accuracy  
## 0.4301639  
  
dtree_function(datax, genres)$overall["Accuracy"]  
  
## Accuracy  
##      0.5
```

## ANALYSIS

As we can see, the accuracy is incredibly poor, even when we picked genres that seem to be different. Let's try to analyze the cause of this by analysing.

### kNN Model

```
genres <-
  list(
    "canadianpop",
    "electronica",
    "rock",
    "modernblues",
    "r&b",
    "polishblackmetal",
    "videogamemusic",
    "irishfolk",
    "koreanpop",
    "hiphop"
  )
knn_function(datax, genres)

## Confusion Matrix and Statistics
##
##           Reference
## Prediction   canadianpop electronica hiphop irishfolk koreanpop
##   canadianpop      355       24     105      36       7
##   electronica        15       36       7      2       2
##   hiphop            103       8     526      41       3
##   irishfolk         24       2     28       7       0
##   koreanpop          3       3       5      6       0
##   modernblues       140      31     99      27       2
##   polishblackmetal     0       1       0      0       0
##   r&b              9       3     18      12       1
##   rock              21      11     18      12       2
##   videogamemusic     37       3     49      14       2
##
##           Reference
## Prediction   modernblues polishblackmetal r&b rock videogamemusic
##   canadianpop      152           0    22    24      40
##   electronica        16           0     2     4      3
##   hiphop            130           1    17    20      52
##   irishfolk          35           0     9     4      12
##   koreanpop          3           0     0     1      0
##   modernblues       311           5    28    32      50
##   polishblackmetal     2           0     1     0      1
##   r&b              25           1    10    0      0
##   rock              43           0     7    36      3
##   videogamemusic     47           0     4     6      31
##
## Overall Statistics
##
##           Accuracy : 0.4302
##           95% CI  : (0.4125, 0.448)
##   No Information Rate : 0.2803
```

```

##      P-Value [Acc > NIR] : < 2.2e-16
##
##          Kappa : 0.2781
##
##  McNemar's Test P-Value : NA
##
## Statistics by Class:
##
##          Class: canadianpop Class: electronica Class: hiphop
## Sensitivity           0.5021        0.29508       0.6152
## Specificity            0.8250        0.98258       0.8292
## Pos Pred Value         0.4641        0.41379       0.5838
## Neg Pred Value         0.8460        0.97098       0.8469
## Prevalence              0.2318        0.04000       0.2803
## Detection Rate          0.1164        0.01180       0.1725
## Detection Prevalence     0.2508        0.02852       0.2954
## Balanced Accuracy        0.6636        0.63883       0.7222
##
##          Class: irishfolk Class: koreanpop Class: modernblues
## Sensitivity            0.044586       0.000000      0.4071
## Specificity             0.960595       0.993072      0.8189
## Pos Pred Value          0.057851       0.000000      0.4290
## Neg Pred Value           0.948788       0.993727      0.8052
## Prevalence                0.051475       0.006230      0.2505
## Detection Rate            0.002295       0.000000      0.1020
## Detection Prevalence      0.039672       0.006885      0.2377
## Balanced Accuracy          0.502590       0.496536      0.6130
##
##          Class: polishblackmetal Class: r&b Class: rock
## Sensitivity            0.000000       0.100000      0.28346
## Specificity             0.998357       0.976610      0.95997
## Pos Pred Value          0.000000       0.126582      0.23529
## Neg Pred Value           0.997701       0.969707      0.96859
## Prevalence                0.002295       0.032787      0.04164
## Detection Rate            0.000000       0.003279      0.01180
## Detection Prevalence      0.001639       0.025902      0.05016
## Balanced Accuracy          0.499178       0.538305      0.62172
##
##          Class: videogamemusic
## Sensitivity             0.16146
## Specificity              0.94332
## Pos Pred Value            0.16062
## Neg Pred Value             0.94365
## Prevalence                 0.06295
## Detection Rate               0.01016
## Detection Prevalence        0.06328
## Balanced Accuracy            0.55239

```

As we can see from the results, several genres are more obscure, and/or have only a small number of songs in the given dataset. Consequently, the KNN algorithm will find very few neighbors of these genres when trying to classify any given point. Therefore, it makes sense that the classification accuracy and other stats would be poor, since KNN makes a decision based on label popularity. If we only have 5 nearby labels (genres) to look at, and each one is different (due to the low proportion of songs in each of the nearby genres), then it's essentially a toss-up for assigning a predicted label.

Let's analyse the **Decision Tree** model:

```

genres <-
  list(
    "canadianpop",
    "electronica",
    "rock",
    "modernblues",
    "r&b",
    "polishblackmetal",
    "videogamemusic",
    "irishfolk",
    "koreanpop",
    "hiphop"
  )
dtree_function(datax, genres)

## Confusion Matrix and Statistics
##
##          Reference
## Prediction   canadianpop electronica hiphop irishfolk koreanpop
##   canadianpop      101        7     31      10       2
##   electronica        0        0      0       0       0
##   hiphop            16        2     111      2       0
##   irishfolk          3        0      0       3       0
##   koreanpop          0        0      0       0       0
##   modernblues        44        1     15      11       1
##   polishblackmetal     0        0      0       0       0
##   r&b              0        0      0       0       0
##   rock              0        0      0       0       0
##   videogamemusic      3       12     17      2       0
##          Reference
## Prediction   modernblues polishblackmetal r&b rock videogamemusic
##   canadianpop      54        0     6     9       0
##   electronica        0        0      0      0       0
##   hiphop            13        0     6     4       1
##   irishfolk          3        0      0      0       3
##   koreanpop          0        0      0      0       0
##   modernblues        64        2     3    16       0
##   polishblackmetal     0        0      0      0       0
##   r&b              0        0      0      0       0
##   rock              0        0      0      0       0
##   videogamemusic      6        0      0      0      26
##
## Overall Statistics
##
##           Accuracy : 0.5
##           95% CI : (0.4596, 0.5404)
##   No Information Rate : 0.2852
##   P-Value [Acc > NIR] : < 2.2e-16
##
##           Kappa : 0.3451
##
##   Mcnemar's Test P-Value : NA

```

```

## Statistics by Class:
##          Class: canadianpop Class: electronica Class: hiphop
## Sensitivity           0.6048        0.00000      0.6379
## Specificity           0.7314        1.00000      0.8991
## Pos Pred Value        0.4591           NaN       0.7161
## Neg Pred Value        0.8308        0.96393      0.8615
## Prevalence             0.2738        0.03607      0.2852
## Detection Rate         0.1656        0.00000      0.1820
## Detection Prevalence   0.3607        0.00000      0.2541
## Balanced Accuracy      0.6681        0.50000      0.7685
##          Class: irishfolk Class: koreanpop Class: modernblues
## Sensitivity            0.107143      0.000000      0.4571
## Specificity            0.984536      1.000000      0.8021
## Pos Pred Value         0.250000           NaN       0.4076
## Neg Pred Value         0.958194      0.995082      0.8322
## Prevalence              0.045902      0.004918      0.2295
## Detection Rate          0.004918      0.000000      0.1049
## Detection Prevalence    0.019672      0.000000      0.2574
## Balanced Accuracy       0.545839      0.500000      0.6296
##          Class: polishblackmetal Class: r&b Class: rock
## Sensitivity            0.000000      0.00000      0.00000
## Specificity            1.000000      1.00000      1.00000
## Pos Pred Value          NaN           NaN           NaN
## Neg Pred Value          0.996721      0.97541      0.95246
## Prevalence              0.003279      0.02459      0.04754
## Detection Rate          0.000000      0.00000      0.00000
## Detection Prevalence    0.000000      0.00000      0.00000
## Balanced Accuracy       0.500000      0.50000      0.50000
##          Class: videogamemusic
## Sensitivity             0.86667
## Specificity             0.93103
## Pos Pred Value          0.39394
## Neg Pred Value          0.99265
## Prevalence               0.04918
## Detection Rate           0.04262
## Detection Prevalence     0.10820
## Balanced Accuracy        0.89885

```

Decision tree classifiers are great for both binary and multi-class problems. They make decisions based on the values (branches) of attributes (leaves / nodes) of the thing being classified, traversing a tree-like structure until reaching a final classification. So in our case, the leaves would be song attributes like “tempo” or “danceability”, and the branches would be the different values each attribute can take. As we saw before, having very few songs in a particular genre is the most likely culprit of our poor results. There’s also the issue of class imbalance, where some genres (like Hip Hop and Canadian pop) have a far greater percentage of songs than other genres do.

## TRYING TO IMPROVE ACCURACY

Let's see what the most popular genres are.

```
descending <- datax %>% group_by(Genre) %>% summarise(n = n()) %>% arrange(desc(n))  
descending
```

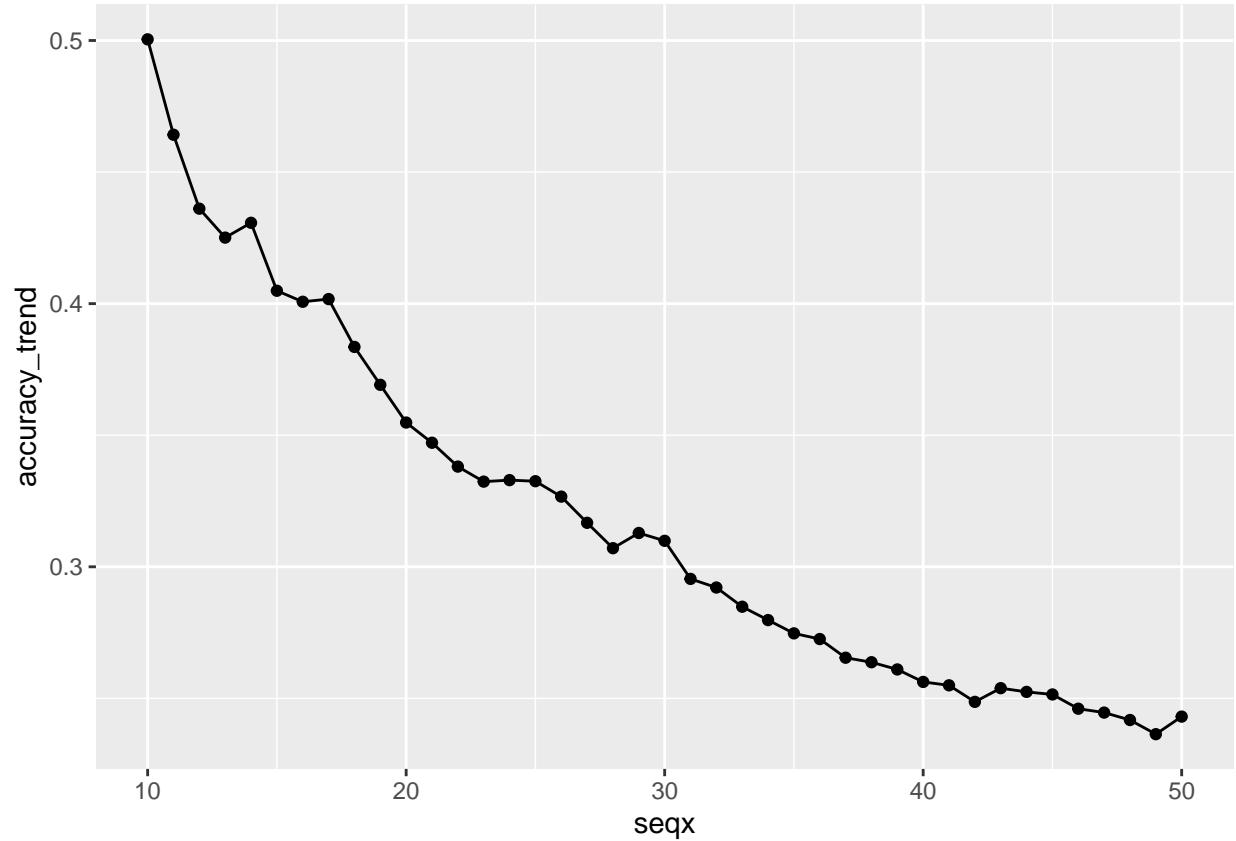
```
## # A tibble: 626 x 2  
##   Genre             n  
##   <fct>            <int>  
## 1 alternativeamericana 1891  
## 2 electrolatino      1009  
## 3 doo-wop            972  
## 4 reading            969  
## 5 nuelectro          909  
## 6 groovemetal        903  
## 7 psychill           901  
## 8 deepdeephouse      892  
## 9 torontoindie       884  
## 10 newrave            875  
## # ... with 616 more rows
```

If we reduce the genres to only those that are very common, the kNN Model in theory should show a considerable improvement as any given label will have plenty of neighbours. Let's try by keeping only the top 30 most popular genres.

```
genre_top30 <- descending$Genre[1:30]  
new_data <- filter(datax, Genre %in% genre_top30)  
  
knn_function(new_data, genre_top30)$overall["Accuracy"]  
  
## Accuracy  
## 0.3098714
```

This accuracy is actually worse than when we hand-picked the genres. Let's see the trend in this accuracy as we pick the top 50 to the top 10 most common genres.

```
seqx <- 50:10  
top_trend <- function(x){  
  genre_topx <- descending$Genre[1:x]  
  new_datax <- filter(datax, Genre %in% genre_topx)  
  knn_function(new_datax, genre_topx)$overall["Accuracy"]  
}  
  
accuracy_trend <- sapply(seqx, top_trend)  
ggplot(data.frame(accuracy_trend), aes(x = seqx, y = accuracy_trend)) + geom_point() + geom_line()
```



Even in our best case scenario, when we use only 10 genres, our accuracy is still only 0.500441

## CONCLUSION

No matter what we try and even when we consider the practically impossible scenario of only including the top 10 most common genres, accuracy does not increase much beyond 0.5. This means that there is massive overlap within the genres. The possible cause for this could be that spotify divides their data into these many genres to aid their machine learning algorithm which predicts user behaviour and not genre. Let's see some overlapping genres:

```
genre_grouped <- datax %>% group_by(Genre) %>%
  summarise(dance = mean(Danceability),
            energy = mean(Energy),
            key = mean(Key),
            loudness = mean(Loudness),
            mode = mean(Mode),
            speechness = mean(Speechiness),
            acousticness = mean(Acousticness),
            intstrumentalness = mean(Instrumentalness),
            liveness = mean(Liveness),
            valence = mean(Valence),
            tempo = mean(Tempo),
            duration = mean(Duration_ms),
            Time = mean(time_signature)) %>%
  arrange(Genre)

genre_grouped

## # A tibble: 626 x 14
##   Genre  dance  energy    key  loudness   mode speechness acousticness
##   <fct> <dbl>  <dbl> <dbl>  <dbl> <dbl>      <dbl>       <dbl>
## 1 ""     0.606  4.42   -7.42   0.577  0.0765    0.463     0.00234
## 2 "abs~" 0.600  0.627   5.31   -10.5    0.5      0.0835    0.247
## 3 "aca~" 0.554  0.578   5.86   -6.88    0.630    0.0807    0.398
## 4 "aco~" 0.540  0.557   4.61   -6.69    0.794    0.0372    0.406
## 5 "afr~" 0.495  0.695   5.19   -6.62    0.893    0.0864    0.356
## 6 "afr~" 0.616  0.668   5.66   -9.00    0.588    0.0660    0.298
## 7 "ala~" 0.521  0.575   4.84   -8.46    0.772    0.0459    0.352
## 8 "alb~" 0.736  0.709   5.70   -6.39    0.394    0.184     0.159
## 9 "alb~" 0.565  0.723   5.51   -5.53    0.919    0.0407    0.190
## 10 "alb~" 0.489  0.621   5.71   -8.74    0.636    0.0556    0.258
## # ... with 616 more rows, and 6 more variables: intstrumentalness <dbl>,
## #   liveness <dbl>, valence <dbl>, tempo <dbl>, duration <dbl>, Time <dbl>
```

Let's see the overlap between Swedish Death Metal and Polish Black Metal:

```
print(genre_grouped[c(which(genre_grouped$Genre == "swedishdeathmetal"),
                           which(genre_grouped$Genre == "polishblackmetal")), ], width = Inf)

## # A tibble: 2 x 14
##   Genre           dance energy   key loudness mode speechness acousticness
##   <fct>        <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
## 1 swedishdeathmetal 0.242  0.935  5.18 -6.62  0.702  0.109   0.00107
## 2 polishblackmetal  0.220  0.868  4.3  -6.32  0.6    0.0985  0.0000567
##   intstrumentalness liveness valence tempo duration Time
##               <dbl>     <dbl>    <dbl>    <dbl>    <dbl> <dbl>
## 1             0.595    0.226   0.221 22908. 256855.  3.95
## 2             0.774    0.227   0.157 26842. 382170   3.4
```

What about Emo Pop and Indonesian Punk Pop

```
print(genre_grouped[c(which(genre_grouped$Genre == "indonesianpoppunk"),
                           which(genre_grouped$Genre == "popemo")), ], width = Inf)

## # A tibble: 2 x 14
##   Genre           dance energy   key loudness mode speechness acousticness
##   <fct>        <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
## 1 indonesianpoppunk 0.472  0.853  4.96 -4.78  0.888  0.0647  0.0647
## 2 popemo          0.478  0.863  4.88 -4.52  0.734  0.0661  0.0288
##   intstrumentalness liveness valence tempo duration Time
##               <dbl>     <dbl>    <dbl>    <dbl>    <dbl> <dbl>
## 1             0.0221   0.200   0.505 26980. 224151.  3.97
## 2             0.0236   0.214   0.512 26510. 206123.  3.91
```

Both of the genres in these sit have striking similarities and these are just a few of the 626 genres that Spotify uses to classify its music data.

Having these varied genres definitely aids in predicting user behaviour, similar to the Netflix Movie Prediction Program we did in the last course. But having these many genres also means that there is going to be major overlap and the distinction between genres is going to be very less. Our current genomes don't have enough range to appropriately house this variety of genres and still produce a desirable result.

What works as a boon when it comes to predicting user behaviour works as a bane when it comes to predicting genres, thus serving as the main limitation.

Perhaps in the future, if the data we received had the name of the artist, our algorithm would perform much better because artists don't usually tend to stray too far away from their genre (other than our occasional Bob Dylan). Also there could be a genre classification system that strikes a balance between user prediction and genre prediction.